

# Debugging Techniques for Distributed R-Trees

Sávio Teles<sup>1</sup>, Jose Ferreira de S. Filho<sup>1</sup>

<sup>1</sup>Instituto de Informática – Universidade Federal de Goiás (UFG)  
Alameda Palmeiras, Quadra D, Câmpus Samambaia  
131 - CEP 74001-970 – Goiânia – GO – Brazil

savioteles@gmail.com, jkairos@gmail.com

**Abstract.** *R-Trees arrived on the scene in 1984 by Antonin Guttman, as a mechanism of handling spatial data efficiently, building an index structure that helps retrieve data items quickly according to their spatial locations. The R-tree has found significant use in both theoretical and applied contexts. However, it brings a big challenge in terms of debugging. In this paper, we argue that distributed debugging is a non-trivial task and few efforts have been developed for distributed debugging of spatial objects based on R-Trees. We use examples of this issue and briefly discuss basic debugging techniques, then we propose RDebug, a technique for debugging a distributed R-Tree.*

## 1. Introduction

The Internet has revolutionized the computer and Geographical Information Systems (GIS) like nothing before. In fact the Internet brought big changes how systems store, retrieve and analyze spatial data. The ever-increasing of the large geospatial datasets and the widely application of the complex geocomputation make the parallel processing of GIS an important component of high-performance computing. Thus, spatial distributed applications came on the scene, a spatial distributed application is one which software components located at networked computers communicate and coordinate their actions for geoprocessing.

In order to handle spatial data efficiently, a database system needs an index mechanism that will help it retrieve data items quickly according to their spatial locations. An R-tree is a dynamic index structure, which meets this requirement and is broadly used in GIS. Generally, the R-Tree index is built and processed in a single machine. However, process an R-Tree in a single machine is not feasible because of the huge size of the geospatial datasets. Thus, many researches such as [An et al. 1999, de Oliveira et al. 2011, Zhong et al. 2012], show that a distributed index structure spanning the workstations can provide an efficient shared storage structure that can be used to gather geographic information more efficiently.

A big challenge though has arisen of spanning the R-Tree index structure among computers: How debug a distributed R-Tree index structure?

Debugging is an essential step in the development process, albeit often neglected in the development of distributed applications due to the fact that distributed systems complicate the already difficult task of debugging [WH Cheung 1990].

In recent years, researchers have been developed some helpful debugging techniques for distributed environment. Nevertheless, we have not found any technique to

debug a distributed R-Tree. In this paper, we propose a debugging technique for data insertion in an R-Tree, besides of debugging the quality of the index built on the cluster.

The debug method, hereafter called RDebug, uses the distributed index structure to aggregate debugging information. RDebug is used by DistGeo, a shared-nothing platform for distributed spatial algorithms processing. We also created a graphical tool to visualize the debugging information and the R-Tree index structure, called RDebug Visualizer.

The main achievements of this paper are:

- RDebug - A debugging technique to the data insertion algorithm in a distributed R-Tree.
- DistGeo - A Shared-nothing platform, without point of failure, to process distributed spatial algorithms of an R-Tree.
- RDebug Visualizer - A graphical tool to visualize debugging information and the distributed R-Tree index.

The rest of the paper is structured as follows. Section 2 describes the distributed processing of spatial algorithms, Section 3 presents our approach for distributed R-Tree debugging. In Section 4, we briefly give an overview of the use of debugging techniques for distributed environments and the view of the distributed spatial algorithms. Finally, we close the paper with some concluding remarks in Section 5.

## 2. Distributed Processing of Spatial Algorithms

### 2.1. Data Structures for Spatial Data Processing

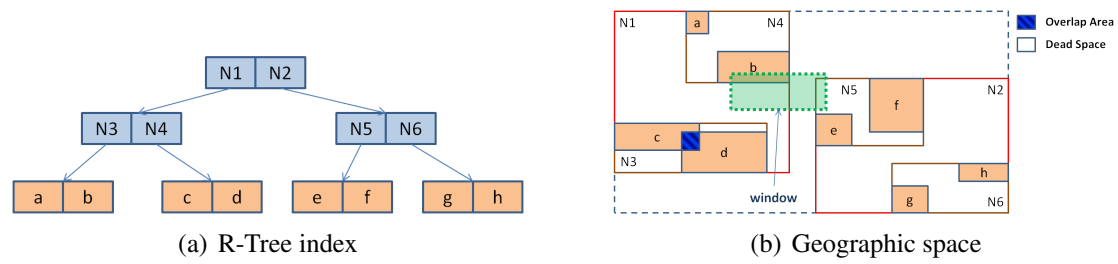
An index based on objects spatial locations is desirable, but classical one-dimensional database indexing structures are not appropriate to multi-dimensional spatial searching structures based on exact matching of values, such as hash tables, are not useful because a range search is required. Structures using one-dimensional ordering of key values, such as B-trees [Bayer and McCreight 1970, Bentley 1975b] and ISAM indexes, do not work because the search space is multi-dimensional.

A number of structures have been proposed for handling multi-dimensional point data, such as: KD-Tree [Bentley 1975a], Hilbert R-Tree [Kamel and Faloutsos 1994] and R-Tree [Guttman 1984].

An R-Tree is a height-balanced tree similar to a B-Tree [Comer 1979] with index records in its leaf nodes containing pointers to data objects. The key idea of the data structure is to group nearby objects and represent them with their minimum bounding rectangle (MBR) in the next higher level of the tree. Since all objects lie within this bounding rectangle, a query that does not intersect the bounding rectangle also cannot intersect any of the contained objects.

Figure 1 illustrates the hierarchical structure of an R-Tree with a root node, internal nodes ( $N1...2 \subset N3...6$ ) and leaves ( $N3...6 \subset a...h$ ). The Figure 1(b) shows MBRs grouping spatial objects of  $a...h$  in sets by their co-location. The Figure 1(a) illustrates the R-Tree representation. Each node stores at most  $M$  and at least  $m \leq M/2$  entries [Guttman 1984].

The Window Query is one of major query algorithms in R-Tree. The input is a search rectangle (Query box). The search starts from the root node of the tree. Every



**Figure 1. R-Tree Structure**

internal node contains a set of rectangles and pointers to the corresponding child node and every leaf node contains the rectangles of spatial objects (the pointer to some spatial object can be there).

For every rectangle in a node, it has to be decided if it overlaps the search rectangle or not. If yes, the corresponding child node has to be searched also. Searching is done in a recursively until all overlapping nodes have been traversed. When a leaf node is reached, the contained bounding boxes (rectangles) are tested against the search rectangle and their objects (if there are any) are put into the result set if they lie within the search rectangle.

In Figure 1, the search starts on root node, where window intersects with the N1 and N2 nodes. Then, the algorithm analyze N4 and n5 entries in N1 node, in which only N4 intersects with the window. Analyzing N4, the algorithm returns the object namely 'b', that is the single object that intersects the window.


In N2, none entry intersects with the window. This happens because of dead space, in other words, the window intersects with a space which not contains any data. The dead space should be minimized to improve performance, since decisions which paths have to be traversed can be taken on higher levels. The overlap area between rectangles should be minimized too because it degrades the performance of R-Tree [Beckmann et al. 1990]. Less overlapping reduces the amount of sub-trees accessed during r-tree traversal. The area between c and d in Figure 1 is an example of overlapping.

## 2.2. DistGeo: A Platform of Distributed Spatial Operations for Geoprocessing

This work developed a platform, namely DistGeo, to process the spatial operations in a cluster of computers. DistGeo is based on the shared-nothing architecture, which the nodes do not share CPU, hard disk and memory and the communication relies on message exchange. Figure 3 depicts DistGeo plataform based on peer-to-peer model, with the data manage by the cluster presented as a ring topology. It is divided in ranges of keys, which are managed for each node of the cluster. To a node join the ring it must first receive a range.

The range of keys are known by each node in the cluster. For instance, in a ring representation, whose key set start with 0 till 100, if we have 4 nodes in the cluster, the division could be done as shown below: a) 0-25, b) 25-50, c) 50-75 e d) 75-100. If we want to search for one object with key 34, we certainly should look on the node 2.

Every replica of an object is equally important, in other words, there is not a master replica. Read/Write operations may be performed in any node of the cluster. When a request is made to a cluster's node, it becomes the coordinator of the operation requested




**Figure 2. DistGeo Platform**

by the client. The coordinator works as a proxy between the client and the cluster's nodes.

DistGeo utilizes the Gossip protocol, which every cluster node exchanges information among themselves for service discovery and knowing the statuses of the cluster's nodes. In the Gossip protocol every second a message is exchanged among three nodes in the cluster, consequently every cluster's node have knowledge of each other.

Figure 3 illustrates the structure of a Distributed R-Tree in a cluster. The partitioning it is performed grouping the nodes in cluster and creating the indexes according to the R-Tree structure. The lines in Figure 3 show the need for message exchange to reach the sub-trees during the algorithm processing.



**Figure 3. R-Tree Partitioning in DistGeo**

Insertions and searching in a distributed R-Tree are similar to the non-distributed version, except for: i) The need of message exchange to access the distributed partitions, ii) Concurrency control and consistency due to the parallel processing in the cluster. Both were implemented in DistGeo.

The distributed index has been built according to the taxonomy defined in [An et al. 1999], as follows: i) Allocation Unit: block - A partition is created for every R-Tree node; ii) Allocation Frequency: overflow - In the insert process, new partitions are created when a node in the tree needs to split; iii) Distribution Policy: balanced - To keep the tree balanced the partitions are distributed among the cluster nodes.

Reliability and fault-tolerance are implemented storing the R-Tree nodes in multiple nodes in the cluster. Each node N receives a key, which is used to store the node in a server S responsible for ring range, replicating the node N to the next two servers in

S (clockwise). In a distributed R-Tree, the requests are always sent to the cluster's node that stores the root node of the R-tree. If a message is sent to a node  $N$ , in the moment of R-Tree traversal an active server is elected, and replicates the data in this node to further requisition processing

Some optimizations proposed in [Beckmann et al. 1990] to reduce the overlap and the dead area were adapted and implemented on DistGeo. Reducing the overlap and dead area minimize the number of messages exchange in network on search algorithms, because the query access less nodes during the tree traversal (see Sub-Section 2.1). This works implements a new distributed debug algorithm of the R-Tree index (RDebug) that helps to reduce the overlap and dead area. This algorithm will be presented in Section 3.

### 3. A Technique for Debugging A Distributed R-Tree

Spatial index debugging is a big challenge in a distributed R-Tree and this section describes a new technique RDebug, which allows debugging the index building of an R-Tree.

The R-tree index building follows a top-down approach, in other words, the index is always built from root to leaves. Debugging the index reliability as the index is built is a non-trivial task, and the aim of this paper is to show a technique for index debugging after it has been created. Common challenges when working with an R-Tree are: i) Reliability of the nodes replicas of the R-Tree, ii) ensure that the MBR of the parent nodes intersect the MBR of their children, iii) the existence of duplicated nodes or being referenced by more than one parent node, and iv) if the value  $M$  and  $m$  of the nodes are compliant with the R-Tree descriptions as shown in Section 2. Furthermore, it is possible to access index data to help in its optimization as dead space and overlapping area.

Algorithm 1 shows the RDebug technique for debugging the distributed spatial index, using the index structure itself. The algorithm has two steps: 1) The algorithm processing is similar to the search in an R-Tree; 2) The algorithm does the inverse of a search in an R-Tree appending information to the distributed index.

The first step, called S1 [Search sub-trees] (lines 1 - 11), the Algorithm 1 traverses every node of the R-Tree starting from the root node to the leaves. Its purpose is to spread the debugging algorithm. The first request is sent to any server, which stores a replica of the root node.

If the node  $T$  is not a leaf (lines 2 - 8), then the number of children entries is stored to control the number of expected answers to this node in the second step of the algorithm. This information is stored in a shared memory accessed by all servers with a replica of  $T$ . Lines 4 -7, show that for every entry  $E$  in the node, a message is sent (continuing step S1) to any server that holds a replica of the child node of  $E$ , carrying on the first step in the children nodes. If the node is a leaf, the second step, named S2 [Aggregation] is started.

Second step aim (lines 12 - 39) is to aggregate the information used for future debugging. This step receives the debugging information of every child node of  $T$ . Therefore, for a given node  $T$  with  $n$  children, the second step is invoked  $n$  times in the node  $T$ .

The index itself is used to aggregate this information, the computational resources of the cluster helps improve the debugging information aggregation time. The index re-

verse structure allows, besides of spanning the aggregation information processing, build the debugging aggregation information, as one node of the R-Tree is responsible to aggregate only the information of its children.

The information of the children nodes are stored in a shared memory, with concurrency control, by the replicas of  $T$ . Hence, line 13, those information are retrieved from the shared memory. Line 14 verifies the consistency of  $T$  in the servers that store any replica of  $T$ . Line 15 verifies the consistency of  $M$  and  $m$  values. Lines 16 and 17 calculate the overlap and the dead space area for each node of the R-tree, respectively. Those information help the insertion algorithm designer analyze the quality of the built index. Line 18 get the MBR of the  $T$ . This data are inserted in *information* on line 18. This information will be used as an input to a tool capable of visualizing the index of the R-Tree.

If the aggregation step is being executed in the leaves (lines 20 - 24), then if  $T$  is the root node (line 22), the node information are sent to the client application. If  $T$  is not the root node, in line 24, the information are sent to the parent node of  $T$ . If the aggregation step is in an internal node (lines 26 - 39), the algorithm aggregates the information of the children nodes. In the line 29, the algorithm receives the information sent by the child node. Line 27, verifies if the MBR of the entry that points to the child node is indeed the same MBR sent by the child node.

Line 28 adds the data processed from lines 26 and 27 in *information*. Line 29 acquires the number of children nodes that not sent debugging information yet. This value is stored in the variable *count*, which is decremented and the values is stored on shared memory to let the other replicas know.

If every node has sent the answer, the variable count then will hold the value 0 and lines 30-35 are processed. If  $T$  is the root node, then the information are sent to the client application, otherwise, those information are sent to the parent node of  $T$ . If the variable *count* is greater than 0, then the client information are stored in the shared memory to be used until every child node send replies.

The algorithm 1 was implemented in the DistGeo platform to collect the debugging information of the built distributed R-tree. Those information are used in the platform to find out indexing issues and to optimize the R-tree index for searching. With the aid of RDebug 1 algorithm, it is possible debug the searching algorithms of an R-Tree. E.g: The Window Query algorithm shown on Section 2.1. To tweak RDebug to Window Query, it is only needed add an window query in the first step and gather the aggregation information of the accessed nodes. Whereas, the algorithms that access diverse R-Trees, such as Spatial Join, need a deep change, as the algorithms can go through different paths.

---

**Algorithm 1:** *RDebug(T)*

---

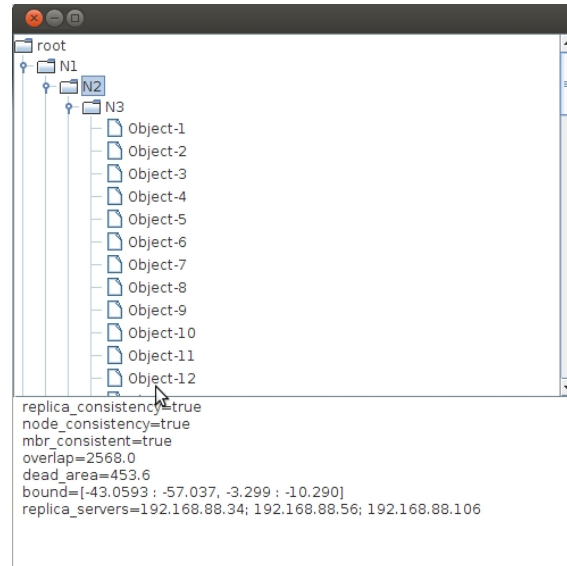
**Data:**  $T$  reference to root node of R-Tree  $tree$

**Result:** Debugging information about distributed R-Tree  $tree$

```
1 S1 [Search subtrees]
2 if  $T$  is not leaf then
3     stores the number of children entries in each replica server of  $T$ 
4     for each entry  $E$  in  $T$  do
5          $server \leftarrow$  choose one server, randomly, that stores one replica of  $E$ 
6         send msg to  $server$  to process the node's child of  $E$  on step S1
7     end
8 else
9     verify the consistency of  $T$  in others replicas
10    Invoke step S2 [Aggregation]
11 end
12 S2 [Aggregation]
13  $information \leftarrow$  the child's information stored on shared memory by
    replicas of  $T$ 
14  $replica\_consistency \leftarrow$  verify the consistency of  $T$  in others replicas
15  $node\_consistency \leftarrow$  verify the consistency of  $M$  and  $m$  values of  $T$ 
16  $overlap \leftarrow$  overlap area of  $T$ 
17  $dead\_area \leftarrow$  dead area of  $T$ 
18  $bound \leftarrow$  MBR of  $T$ 
19 add in  $information$ :  $replica\_consistency$ ,  $node\_consistency$ ,  $overlap$ ,
     $dead\_area$ ,  $bound$ 
20 if  $T$  is leaf then
21     if  $T$  is root then
22         send response with R-Tree nodes information to app client
23     end
24     send msg with  $information$  to parent of  $T$ 
25 else
26      $entry\_info \leftarrow$  information sent by child node
27      $mbr\_consistent \leftarrow$  verify if the bound of the child node is equal to
        bound of entry of  $T$  that points to this child
28     add in  $information$ :  $entries\_info$ , and  $mbr\_consistent$ 
29      $count \leftarrow$  retrieve the number of child entries, which did not send a
        debugging response and decrement by 1
30     if  $count == 0$  then
31         if  $T$  is root then
32             send response with  $information$  to app client
33         else
34             send msg with  $information$  to parent of  $T$ 
35         end
36     else
37         store  $information$  on shared memory
38     end
39 end
```

---

Figure 4 shows a graphical tool (RDebug Visualizer) created to visualize the structure of the distributed R-Tree index, using as the input the information generated by the distributed debugging algorithm in DistGeo platform. The RDebug Visualizer allows analyze the debug information in each node of the R-Tree. The output of the RDebug algorithm shows which nodes are currently inconsistent, so the user can access the path of the node and visualize the node's inconsistent information.



**Figure 4. RDebug Visualizer**

## 4. Related work

Several techniques exist for debugging distributed, concurrent, and parallel programs. Although none of the existing approaches provide support for debugging a distributed R-Tree index [Manolopoulos et al. 2003, Jacox and Samet 2007]. Along with our discussion, we highlight the techniques that influenced the design of RDebug, a new algorithm for debugging a distributed R-Tree index and a peer-to-peer platform to process distributed spatial algorithms.

Researches on distributed spatial data either show techniques to debug distributed applications in general or techniques for R-tree distributed processing. The Section 4.1 shows the distributed debugging researches and 4.2 describes researches of platforms for processing distributed spatial algorithms.

### 4.1. Distributed Debugging Techniques

In [G. et al. 2011] the author breaks down debuggers in two main families: log-based debuggers (also known as post-mortem debuggers) and breakpoint-based debuggers (also known as online debuggers). Log-based debuggers insert log statements in the code of the program to be able to generate a trace log during its execution. Breakpoint-based debuggers, on the other hand, execute the program in the debug mode that allows programs to pause/resume the program execution at certain points, inspect program state, and perform step-by-step execution.



Several breakpoint-based debuggers have been designed for parallel programs using message passing communication including p2d2 [Hood 1996], TotalView [Gottbrath 2009], and Amoeba [Elshoff 1989]. These debuggers offer the traditional commands to, e.g. stop, inspect and step-by-step execution of a running program. Some of them allow to set breakpoints on statements of one process (e.g. TotalView) or a set of processes (e.g. p2d2, prism).

An interesting alternative to traditional breakpoints is message breakpoints {debuggingMP1997}. A message breakpoint stops all receiver processes of the next message sent by a process. The notion of message-breakpoints maps well to the event-loop concurrency model, since it allows to define sensible stepping semantics at message passing level.

A great body of concurrent and parallel debugging techniques are event-based. Event-based debuggers [C. E. McDowell 1989] conceive the execution of a program as a sequence of events. The debugger records the history of the events generated by the application, which can then be used to either browse the events once the application is finished [Fonseca et al. 2007, Stanley et al. 2009], or to replay the execution to recreate the conditions under which the bug was observed. Event-based debuggers have been mainly criticized because the recording process is costly and browsing event history does not scale since manually inspecting huge traces becomes cumbersome and difficult.

[WH Cheung 1990] describes a process for distributed debugging in general and does not focus on a specific debugger or a particular technique, the paper focus is on defining a step by step approach to tackle distributed debugging independent of the environment.

## **4.2. Distributed Spatial Algorithms**

This Section describes briefly the researches that present the use of parallelism in order to improve the response time of the spatial algorithms. M-RTree [Koudas et al. 1996] was the first published paper, which shows a shared-nothing architecture, with a master and several workstations connected to a LAN. The master machine handles a high volume of computation, besides of processing some directories of the R-Tree, it merges the answers to the client machines. Same technique is found on MC-RTree [Schnitzer and Leutenegger 1999].

In [An et al. 1999], it is described that a Network of Workstations (NOW) - also referred to as cluster systems - is a cost-effective solution for high performance. The paper proposes an architecture similar to M-RTree and MC-RTree, with the same disadvantages of both.

In SD-RTree [du Mouza et al. 2007], a binary tree is used instead of an R-Tree. The binary tree increases the number of messages, since the data representation in a binary tree has more levels than the same data represented in an R-Tree. Hadoop-GIS [Kerr 2009] shows a scalable and high performance spatial data warehousing system for running large scale spatial queries on Hadoop. However the gain running large scale queries, it does not use indexing to improve the performance of the operations in the datasets.

[de Oliveira et al. 2011] presents a distributed platform for spatial operations. Al-

though, the solution proposed implements a distributed index, it is not scalable, since every message go through the replicated master node. [saviosbrc] shows a hybrid peer-to-peer platform to process the distributed spatial joint. The architecture presented in [saviosbrc] comprehends a set of machines for naming resolution and events management. Thus, these machines could be a bottleneck in the system.

[Xie et al. 2008] introduces a two-phase load-balancing scheme for the parallel GIS operations in distributed environment.

[Zhang et al. 2009] describes MapReduce and shows how spatial queries can be naturally expressed in this model, without explicitly addressing any of the details of parallelization. Although the geocomputation high performance with this approach, it is only indicated for non-indexed datasets.

In [Zhong et al. 2012], an approach is proposed for "indexing + MapReduce" data processing architecture to improve the computation capability of spatial query .The spatial index is distributed in two levels: locally and globally.

A number of techniques and platforms have been proposed for handling big spatial data, nevertheless none of them propose a platform using a peer-to-peer approach for processing distributed spatial algorithms as found on DistGeo platform. Besides, none of the researches propose a technique for distributed spatial index debugging of an R-Tree.

## 5. Conclusion

DistGeo platform presents an approach for processing the distributed spatial operations through the distributed R-Tree index. Due to the distributed processing nature on this platform an issue arises: Debugging the insertion algorithm when the R-tree nodes a distributed in a cluster.

We have seen researches on spatial data processing and distributed debugging, but none of them propose techniques for debugging the spatial algorithms of an R-Tree. In this paper, we present RDebug, a technique for debugging the distributed index building of an R-Tree. RDebug, uses the R-Tree index itself to gather the debug information. The information gathering is done in the R-Tree index using a down-top traversal in the tree. Utilizing the distributed index itself, the data gathering can be achieve in a distributed way, improving the debugging algorithm efficiency.

RDebug, has been implemented in DistGeo platform. The R-Tree nodes are distributed and replicated over the cluster. Thus, RDebug can be processed without bottlenecks and point of failures. Besides, the R-Tree replicated nodes in the cluster allow load-balancing in the distributed R-Tree index traversal. During the traversal, at every node access of the R-Tree, the traversal might go to a node of the cluster with less workload increasing RDebug performance. The information exchange of the machines statuses is done trough the Gossip algorithm.

A graphical tool(RDebug Visualizer) has been created to visualize the structure of the distributed R-Tree index and the debugging information about the index building. With these input we can identify discrepancies in the index building and optimize it.

Ongoing work includes modify the RDebug algorithm to debug the Window Query and Join Query searching algorithms. The RDebug algorithm is easily adapted

to gather debugging information for Window Query. Whereas, for Join Query algorithm, RDebug must be changed considerably, since the traversal is processed in two different distributed R-Trees. Another ongoing work is to realize performance tests of the RDebug algorithm in DistGeo platform.

## References

- An, N., Lu, R., Qian, L., Sivasubramaniam, A., and Keefe, T. (1999). Storing spatial data on a network of workstations. *Cluster Computing*, 2(4):259–270.
- Bayer, R. and McCreight, E. (1970). Organization and Maintenance of Large Ordered Indices. *ACM-SIGFIDE Workshop on Data Description and Access*, pages 107–141.
- Beckmann, N., Kriegel, H., Schneider, R., and Seeger, B. (1990). *The R\*-tree: an efficient and robust access method for points and rectangles*, volume 19. ACM.
- Bentley, J. (1975a). Multidimensional binary search trees used for associative searching. *Communications of the ACM*, 18(9):517.
- Bentley, J. L. (1975b). Multidimensional Binary Search Trees Used for Associative Searching. *ACM*, pages 509–517.
- C. E. Mcdowell, D. P. H. (1989). Debugging concurrent programs. *ACM Computing Surveys*, 21:593–622.
- Comer, D. (1979). Ubiquitous B-tree. *ACM Computing Surveys (CSUR)*, 11(2):121–137.
- de Oliveira, T., Sacramento, V., Oliveira, S., Albuquerque, P., Cardoso, M., Bloco, I., and Campus, I. (2011). DSI-Rtree - Um Índice R-Tree Escalável Distribuído. In *XXIX Simpósio Brasileiro de Redes de Computadores e Sistemas Distribuídos*.
- du Mouza, C., Litwin, W., and Rigaux, P. (2007). Sd-rtree: A scalable distributed rtree. In *Data Engineering, 2007. ICDE 2007. IEEE 23rd International Conference on*, pages 296–305. IEEE.
- Elshoff, I. J. P. (1989). A distributed debugger for amoeba. In *In Symposium on Parallel and distributed tools*, pages 1–10. SIGPLAN Not.
- Fonseca, R., Porter, G., Katz, R. H., Shenker, S., and Stoica, I. (2007). X-Trace: A pervasive network tracing framework. In *4th USENIX Symposium on Networked Systems Design Implementation*, pages 271–284. Cambridge MA, USA.
- G., B. E., V., C. T., C., N., D., M. W., and D, H. T. (2011). REME-D: a Reflective Epidemic Message-Oriented Debugger for Ambient-Oriented Applications. *ACM*, pages 1275–1281.
- Gottbrath, C. (2009). Deterministically troubleshooting network applications. In *Technical report, TotalView Technologies*. TotalView Technologies.
- Guttman, A. (1984). *R-trees: a dynamic index structure for spatial searching*, volume 14. ACM.
- Hood, R. (1996). The p2d2 project: building a portable distributed debugger. In *In Symposium on Parallel and distributed tools*, pages 127–136. ACM.
- Jacox, E. and Samet, H. (2007). Spatial join techniques. *ACM Transactions on Database Systems (TODS)*, 32(1):7.

- Kamel, I. and Faloutsos, C. (1994). Hilbert R-tree: An Improved R-tree using Fractals. In *VLDB 20th*, page 509. Morgan Kaufmann Publishers Inc.
- Kerr, N. (2009). Alternative Approaches to Parallel GIS Processing. *Arizona State University - Master Thesis*.
- Koudas, N., Faloutsos, C., and Kamel, I. (1996). Declustering spatial databases on a multi-computer architecture. *Advances in Database Technology-EDBT'96*, pages 592–614.
- Manolopoulos, Y., Nanopoulos, A., Papadopoulos, A., and Theodoridis, Y. (2003). R-trees have grown everywhere. *Submitted to ACM Computing Surveys*.
- Schnitzer, B. and Leutenegger, S. (1999). Master-client r-trees: A new parallel r-tree architecture. In *Scientific and Statistical Database Management, 1999. Eleventh International Conference on*, pages 68–77. IEEE.
- Stanley, T., Close, T., and Miller, M. (2009). Causeway: A message-oriented distributed debugger. In *Technical Report HPL-2009-78*. HP Laboratories.
- WH Cheung, JP Black, E. M. (1990). A Framework for Distributed Debugging. *IEEE*, pages 106–115.
- Xie, Z., Ye, Z., and Wu, L. (2008). A two-phase load-balancing framework of parallel gis operations. In *Geoscience and Remote Sensing Symposium, 2008. IGARSS 2008. IEEE International*, volume 2, pages II–1286. IEEE.
- Zhang, S., Han, J., Liu, Z., Wang, K., and Feng, S. (2009). Spatial queries evaluation with mapreduce. In *Grid and Cooperative Computing, 2009. GCC'09. Eighth International Conference on*, pages 287–292. IEEE.
- Zhong, Y., Han, J., Zhang, T., Li, Z., Fang, J., and Chen, G. (2012). Towards parallel spatial query processing for big spatial data. In *Parallel and Distributed Processing Symposium Workshops & PhD Forum (IPDPSW), 2012 IEEE 26th International*, pages 2085–2094. IEEE.